

deckard: A Declarative Tool for Machine Learning Robustness Evaluations

15 April 2025

Summary

The software package presented, called **deckard**, is a modular software toolkit designed to streamline and standardize experimentation in machine learning (ML) with a particular focus on the adversarial scenario. It provides a flexible, extensible framework for defining, executing, and analysing end-to-end ML pipelines in the context of a malicious actor. As it is built on top of the Hydra configuration system, **deckard** supports declarative YAML-based configuration of data preprocessing, model training, and adversarial attack pipelines, enabling reproducible, framework-agnostic experimentation across diverse ML settings.

In addition to configuration management, **deckard** includes a suite of utilities for distributed and parallel execution, automated hyperparameter optimisation, visualisation, and result aggregation. The tooling abstracts away much of the engineering overhead typically involved in adversarial ML research, allowing researchers to focus on algorithmic insights rather than implementation details. The presented software facilitates rigorous benchmarking by maintaining an auditable trace of configurations, random seeds, and intermediate outputs throughout the experimental lifecycle.

The system is compatible with a variety of ML frameworks and several classes of adversarial attacks, making it a suitable backend for both large-scale automated testing and fine-grained empirical analysis. By providing a unified interface for experimental control, **deckard** accelerates the development and evaluation of robust models, and helps close the gap between research prototypes and verifiable, reproducible results.

Statement of need

While tools such as **mlflow** (Zaharia et al. 2018), **Weights & Biases** (Biewald 2020), **optuna** (Akiba et al. 2019), and **Kubernetes** (Kubernetes 2019) provide essential infrastructure for model tracking and experiment management, **deckard** occupies a different position in the ML ecosystem—focusing specifically on configurable, adversarially robust experimentation.

Unlike **MLflow** and **Weights & Biases**, which emphasize logging, visualization, and reproducibility for various ML frameworks, **deckard** enforces reproducibility by construction through its declarative, YAML-driven configuration system built on Facebook’s **hydra** (Yadan 2019) configuration management tool. In contrast to cloud-management software like **Kubernetes**—which is a general-purpose container orchestration platform—**deckard** abstracts away orchestration details and offers native support for parallel and distributed experimentation, tailored to ML workflows involving attack/defense cycles, model retraining, or optimisation. While **deckard** integrates tightly with IBM’s **Adversarial Robustness Toolbox** (Nicolae et al. 2018), the software is designed to be easily extensible to other attack frameworks. The human- and machine-readable parameter configuration system allows researchers to declaratively define end-to-end pipelines that span data sampling, preprocessing, model training, attack generation, defense evaluation, multi-objective optimisation, and visualisation. Tools like **ray** (Moritz et al. 2018), **optuna** (Akiba et al. 2019), or **nevergrad** (Bennet et al. 2021) offer components of this pipeline (*e.g.*, hyperparameter search or configuration management), but lack unified support for adversarial ML, verification, or auditability at scale. While **deckard** complements these existing tools, and in many cases can be integrated with them, its primary contribution is in automating and verifying adversarial ML experiments in a way that is both extensible and framework-agnostic.

Usage

Various versions of this software have been used in several recently published and not-yet-published works by the author of this paper, all of which are available in the **examples** folder in the source code repository <https://github.com/simplymathematics/deckard>. One published work, now reproducible via the **examples/attack_defence_survey** folder, includes a large survey of attacks and defences against canonical datasets and models (C. Meyers, Löfstedt, and Elmroth 2023). Another work analysed the run-time requirements of attacks against a particular model before and after retraining against those attacks (Charles Meyers et al. 2024) (reproducible via **examples/retraining**). The next paper formalised a method for estimating the time-to-failure of a given model against a suite of attacks and introduce a metric that quantifies the ratio of attack and training cost (Meyers et al. 2023) (reproducible via **examples/survival_heuristic**). Furthermore, a not yet published work uses this time-to-failure model as a mechanism for analysing the cost efficacy of various hardware choices in the context of adversarial attacks (reproducible via **examples/power**) (C. Meyers et al. 2024). Another work exploits the tooling to train a custom model that is designed to run client-side by using compression algorithms to measure the distance between text (reproducible via **examples/compression**).

Experiment Management

Typically ML projects are composed of long and complex pipelines that are highly dependent on a number of parameters that must be configured by either the model builder or attacker. Due to the large scale and cost associated with training ML models, it is often necessary to tune a model using many individual model configurations (often called *hyper-parameters*). To determine adversarial robustness, one of many benchmark datasets is first sampled, then preprocessed, sent to a model, with optional pre- and post-processing defences, and then scored according to some chosen metric which may include the performance against any number of adversarial attacks. Each stage in this example pipeline might include tens or hundreds of possible sets of hyper-parameters that must be exhaustively tested. Furthermore, this problem scales drastically as we include more and more stages in a pipeline since each additional stage introduces a new combinatorial layer of complexity, rapidly expanding the total number of potential configurations that must be evaluated for robustness and be reproducible for posterity. Not only does **deckard** provide a standard way to document and configure these hyper-parameters, it gives each experiment an auditable identifier.

Reproducibility and Auditability

For ML, various regulatory and legal frameworks govern safety (The Parliament of the European Union 2024; “ISO 26262-1:2011, Road Vehicles — Functional Safety” 2018; *IEC 61508 Safety and Functional Safety* 2010; *IEC 62304 Medical Device Software - Software Life Cycle Processes* 2006), privacy (The Parliament of the European Union 2024; European Parliament and Council of the European Union 2016; Legislature of the United States 1996, 1998) and/or transparency (The Parliament of the European Union 2024; The Legislature of California 2024). The software package presented here provides a machine- and human-readable format for creating reproducible and auditable experiments as required by various regulations. In addition, several examples connected to both published and not-yet-published work live in the **examples** folder in the repository, allowing for easy reproducibility of several extensive sets of experiments across several popular ML software frameworks. The **power** example provides a reproducible way to run a suite of adversarial tests using popular cloud-based platforms and the **retraining** and **survival_heuristic** examples provide examples of both CPU and GPU-based parallelisation, respectively.

The **basics** subfolder provides a minimum working example for each of the supported ML frameworks: **tensorflow** (Abadi et al. 2015), **pytorch** (Paszke et al. 1912), **scikit-learn** (Pedregosa et al. 2011), and **keras** (Chollet 2015). The **basics** folder also provides examples of various classes of adversarial examples: *poisoning* attacks that change model behaviour by injecting data during training (Biggio, Nelson, and Laskov 2012), *inference* attacks (Li and Zhang 2021) that attempt to reverse engineer properties of the training data, *extraction* attacks that attempt to reverse engineer the model (Jagielski et al. 2020), and *evasion* attacks that

induce errors of classification during run-time (C. Meyers, Löfstedt, and Elmroth 2023). The parameters file for each experiment ensures that a given pipeline can be reproduced and the standardised format allows us to derive a hash value that is hard to forge but easy to verify. Not only does this hash serve as an identifier to track the state of an experiment, but also serves as a way to audit the parameters file for tampering. Likewise, by using `dvc` (DVC Authors 2023) to track any input or output files specified in the parameters file, the software associates each score file with a identifier that is easy to track and verify, but hard to forge—ensuring that forged or modified results are easy to spot in version-controlled experiment repository.

Parallel and Distributed Design

Since ML projects can exploit specialized hardware such as multi-core processors or GPUs, and often rely on clusters of machines for large-scale data processing, it was necessary to enable parallel and distributed experiment execution and model optimization. By leveraging the `hydra` configuration framework, `deckard` automatically supports optimization libraries like `nevergrad` (Bennet et al. 2021), `Adaptive Experimentation` (A. Developers 2025), and `optuna` (Akiba et al. 2019), making the software modular and extensible. Additionally, experiments can be managed using a variety of popular job schedulers, including `Ray` (Moritz et al. 2018), `Redis Queue` (Stamps 2025), and `slurm` (Yoo, Jette, and Grondona 2003) for distributed jobs or `joblib` (Joblib Developers 2025) for jobs on a single machine.

By using a declarative design, a given set of experiments can be specified once and executed seamlessly across different back-ends without modification to the underlying codebase. This makes `deckard` both adaptable and scalable, suitable for use on personal laptops, multi-node servers, or large-scale, high-performance clusters. When configured appropriately, experiment batches can be parallelized, enabling massive parameter sweeps, ensemble evaluations, or adversarial robustness tests to be executed in parallel—reducing turnaround time while maintaining strong guarantees on reproducibility and auditability. The design of the presented software prioritizes clarity and maintainability by capturing each experimental configuration as a YAML artifact, making both successful and failed runs equally traceable and shareable. This approach transforms experiment tracking from an afterthought into a first-class component of the trustworthy ML workflow.

Funding

Financial support has been provided in part by the Knut and Alice Wallenberg Foundation grant number 2019.0352 and by the eSSENCE Programme under the Swedish Government’s Strategic Research Initiative.

Acknowledgements

The author would like to thank Aaron MacSween, Abel Souza, and Mohammad Saledghpour Reza for their guidance in software design principles. In particular, the author appreciates Mohammad’s code and documentation regarding cloud-based and other Kubernetes deployments. The author would like to thanks his advisors, Erik Elmroth and Tommy Löfstedt for their research expertise, funding, and patience.

References

- Abadi, Martín, Ashish Agarwal, Paul Barham, Eugene Brevdo, Zhifeng Chen, Craig Citro, Greg S. Corrado, et al. 2015. “TensorFlow: Large-Scale Machine Learning on Heterogeneous Systems.”
- Akiba, Takuya, Shotaro Sano, Toshihiko Yanase, Takeru Ohta, and Masanori Koyama. 2019. “Optuna: A Next-Generation Hyperparameter Optimization Framework.” In *Proceedings of the 25th ACM SIGKDD*, 2623–31.
- Bennet, Pauline, Carola Doerr, Antoine Moreau, Jeremy Rapin, Fabien Teytaud, and Olivier Teytaud. 2021. “Nevergrad: Black-Box Optimization Platform.” *ACM SIGEVOlution* 14 (1): 8–15.
- Biewald, Lukas. 2020. “Experiment Tracking with Weights and Biases.” <https://www.wandb.com/>.
- Biggio, Battista, Blaine Nelson, and Pavel Laskov. 2012. “Poisoning Attacks Against Support Vector Machines.” *International Conference on Machine Learning*.

- Charles Meyers, Tommy Löfstedt, and Erik Elmroth. 2024. “Massively Parallel Evasion Attacks and the Pitfalls of Adversarial Retraining.” *EAI Endorsed Transactions on Internet of Things* 10.
- Chollet, François. 2015. “Keras.” *GitHub Repository*. <https://github.com/fchollet/keras>; GitHub.
- Developers, Ax. 2025. “AX: Adaptive Experimentation Platform.” <https://ax.readthedocs.io/en/stable/#>.
- Developers, Joblib. 2025. “Joblib: Running Python Functions as Pipeline Jobs.” <https://joblib.readthedocs.io/en/stable/>.
- DVC Authors. 2023. “DVC–Data Version Control.” Github. <https://github.com/iterative/dvc.org>.
- European Parliament, and Council of the European Union. 2016. “Regulation (EU) 2016/679 of the European Parliament and of the Council.” May 4, 2016. <https://data.europa.eu/eli/reg/2016/679/oj>.
- IEC 61508 *Safety and Functional Safety*. 2010. 2nd ed. International Electrotechnical Commission.
- IEC 62304 *Medical Device Software - Software Life Cycle Processes*. 2006. 2nd ed. International Electrotechnical Commission.
- “ISO 26262-1:2011, Road Vehicles — Functional Safety.” 2018. <https://www.iso.org/standard/43464.html> (visited 2022-04-20).
- Jagielski, Matthew, Nicholas Carlini, David Berthelot, Alex Kurakin, and Nicolas Papernot. 2020. “High Accuracy and High Fidelity Extraction of Neural Networks.” In *29th USENIX Security Symposium (USENIX Security 20)*, 1345–62.
- Kubernetes. 2019. “Kubernetes—an Open Source System for Managing Containerized Applications.” Github. <https://github.com/kubernetes/kubernetes>.
- Legislature of the United States. 1996. “Health Insurance Portability and Accountability Act.”
- . 1998. “Children’s Online Privacy Protection Act.”
- Li, Zheng, and Yang Zhang. 2021. “Membership Leakage in Label-Only Exposures.” In *Proceedings of the 2021 ACM SIGSAC Conference on Computer and Communications Security*, 880–95.
- Meyers, Charles, Tommy Löfstedt, and Erik Elmroth. 2023. “Safety-Critical Computer Vision: An Empirical Survey of Adversarial Evasion Attacks and Defenses on Computer Vision Systems.” *Artificial Intelligence Review* 56 (Suppl 1): 217–51.
- Meyers, Charles, Mohammad Reza Saleh Sedghpour, Tommy Löfstedt, and Erik Elmroth. 2024. “A Training Rate and Survival Heuristic for Inference and Robustness Evaluation (TRASHFIRE).” <https://arxiv.org/abs/2401.13751>.
- Meyers, Reza, Löfstedt, and Elmroth. 2023. “A Systematic Approach to Robustness Modelling.” *Springer Artificial Intelligence Review*.
- Moritz, Philipp, Robert Nishihara, Stephanie Wang, Alexey Tumanov, Richard Liaw, Eric Liang, Melih Elibol, et al. 2018. “Ray: A Distributed Framework for Emerging {AI} Applications.” In *13th USENIX Symposium on Operating Systems Design and Implementation (OSDI 18)*, 561–77.
- Nicolae, Maria-Irina, Mathieu Sinn, Minh Ngoc Tran, Beat Buesser, Ambrish Rawat, Martin Wistuba, Valentina Zantedeschi, et al. 2018. “Adversarial Robustness Toolbox V1. 0.0.” *arXiv Preprint arXiv:1807.01069*.
- Paszke, Adam, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, et al. 2019. “Pytorch: An Imperative Style, High-Performance Deep Learning Library.” *arXiv Preprint arXiv:1912.01703* 10.
- Pedregosa, F., G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, et al. 2011. “Scikit-Learn: Machine Learning in Python.” *Journal of Machine Learning Research* 12: 2825–30.
- Stamps. 2025. “RQ: Easy Job Queues for Python.” Stamps, an Indonesian CRM company. <https://python-rq.org/>.
- The Legislature of California. 2024. “AB-2013 Generative Artificial Intelligence: Training Data Transparency.”
- The Parliament of the European Union. 2024. “High-Level Summary of the AI Act.”
- Yadan, Omry. 2019. “Hydra – a Framework for Elegantly Configuring Complex Applications.” Github. <https://github.com/facebookresearch/hydra>.
- Yoo, Andy B, Morris A Jette, and Mark Grondona. 2003. “Slurm: Simple Linux Utility for Resource Management.” In *Workshop on Job Scheduling Strategies for Parallel Processing*, 44–60. Springer.
- Zaharia, Matei, Andrew Chen, Aaron Davidson, Ali Ghodsi, Sue Ann Hong, Andy Konwinski, Siddharth Murching, et al. 2018. “Accelerating the Machine Learning Lifecycle with MLflow.” *IEEE Data Eng. Bull.* 41 (4): 39–45.